

CINDERFALL LITE

A simultaneous-turn strategy game where players read an adaptive rival and chase an objective directed by AI.

2 SECRET COMMANDS

4 POINTS TO WIN

2 BOUNDED AIs

The player problem: **AI opponents are usually invisible**

What players experience

- AI difficulty often feels like hidden bonuses
- Procedural events can invalidate careful plans
- Black-box behaviour is hard to learn or outplay
- More complexity hides rather than proves intelligence



Our product insight

Designed for players who want depth without homework

01

Strategy-curious players

Players attracted to tactical decisions but discouraged by long rules, opaque combat maths, or hour-long matches.

NEED • MEANINGFUL CHOICES IN MINUTES

02

Competitive repeat players

Players who value prediction, adaptation, rematches, and opponents whose behaviour can be studied and countered.

NEED • FAIR INFORMATION TO OUTPLAY

03

Game teams and evaluators

Designers and stakeholders who need AI behaviour to be visible, testable, reliable, and safe enough for a live product.

NEED • DEMONSTRABLE AI VALUE


Three commands. Choose exactly two.

EXPAND



Claim adjacent neutral land. Two Expands can chain into a rapid advance.

RAID



Directly capture adjacent enemy land. The outcome is immediate and visible.

Simple rules. Non-trivial choices.

1 RAID

—

1 GUARD

=

DEFENDED

2 RAIDS

—

1 GUARD

=

CAPTURED



Pressure the north? Race along the south? Or commit both commands to break a Guard?

A credible but incomplete rival intent

PUBLIC • 2-TURN INTENT


RAID • NORTH

The player knows at least one Saltkin command will follow this intent—but not its exact target or second command.

S1 • RAID	NORTH
S2 • EXPAND	SOUTH
S3 • DEFEND	BEACON

The model selects one trusted ID. Deterministic code produces the two legal commands.

The world does not punish players. It changes **where the next conflict matters.**



Cinder selects from **three fair candidates**. It cannot alter terrain, scoring, combat, or the announced coordinate.

The model proposes direction. **Code retains authority.**

Resolved public state

No current secret queue
No identity data



LLM selects a candidate ID

Intent or Beacon location
Evidence keys + short explanation



A complete end-to-end live loop

01

Read

See Saltkin's intent and Cinder's exact next Beacon.



02

Predict

Decide whether the rival will focus, split, or trade tempo with Guard.



03

Seal

Submit two secret commands while every model call remains bounded.



04

Resolve

Simultaneous deterministic results, scoring, movement, audit and fallback.

Success means players can read, predict, and outplay AI

30s

Time to understand

A four-scene square tutorial teaches the complete command vocabulary before the first decision.

6–10m

Target session

Four Beacon points and a ten-turn cap create a complete strategic arc suited to demos and rematches.

100%

Explainable outcomes

Every command result, model selection, exception, and fallback has a visible deterministic explanation.

PLAYTEST KPIs • tutorial comprehension • intent-informed decisions • rematch rate • Guard usage • Beacon/capital win mix • side win-rate gap • fallback completion rate

Our claim:

Compelling AI play does not require infinite generation. It can emerge from intelligence players can **understand, predict, exploit, and outplay**.

AI as readable game state

Intent is not flavour text—it changes the player's command allocation.

The map as a bounded actor

Cinder redirects conflict without invalidating agency.

Trust as part of the experience

Exact previews, deterministic authority, and visible fallback make AI judgeable.

